

```
In [1]: #Processamento e Análise em Big Data
#Tarefa 05
#Grupo 03
#20/04/2021
```

```
In [2]: import pandas as pd
from pyspark.sql import functions as F
from pyspark.ml.feature import PCA
from pyspark.ml.feature import StringIndexer, VectorAssembler
from pyspark.ml import Pipeline
from pyspark.sql.functions import col
```

## Análise Descritiva

```
In [56]: data_ = spark.read.csv('pima_indian_diabetes.csv', header=True)
```

```
In [4]: data_.limit(5).toPandas()
```

Out[4]:

	n_pregnant	plasma_glucose_concentration	blood_pressure	triceps_skinfold_thickness	2hours
0	6	148	72	35	
1	1	85	66	29	
2	8	183	64	0	
3	1	89	66	23	
4	0	137	40	35	

### Checando o Schema do Dataframe

```
In [5]: data_.printSchema()
```

```
root
 |-- n_pregnant: string (nullable = true)
 |-- plasma_glucose_concentration: string (nullable = true)
 |-- blood_pressure: string (nullable = true)
 |-- triceps_skinfold_thickness: string (nullable = true)
 |-- 2hours_serum_insulin: string (nullable = true)
 |-- body_mass_index: string (nullable = true)
 |-- diabetes_pedigree_function: string (nullable = true)
 |-- age: string (nullable = true)
 |-- class: string (nullable = true)
```

Vamos converter os valores string para int e float

```
In [57]: num_cols = ['n_pregnant', 'class', 'age']
for col in num_cols:
    data_ = data_.withColumn(col, data_[col].cast('int'))

num_cols = ['plasma_glucose_concentration', 'blood_pressure', 'triceps_skinfold_
thickness',
            '2hours_serum_insulin', 'body_mass_index', 'diabetes_pedigree_functi
on' ]
for col in num_cols:
    data_ = data_.withColumn(col, data_[col].cast('float'))

data_.printSchema()

root
|-- n_pregnant: integer (nullable = true)
|-- plasma_glucose_concentration: float (nullable = true)
|-- blood_pressure: float (nullable = true)
|-- triceps_skinfold_thickness: float (nullable = true)
|-- 2hours_serum_insulin: float (nullable = true)
|-- body_mass_index: float (nullable = true)
|-- diabetes_pedigree_function: float (nullable = true)
|-- age: integer (nullable = true)
|-- class: integer (nullable = true)
```

**copiando Dataframe**

```
In [7]: schema = data_.schema
X_pd = data_.toPandas()
data_model = spark.createDataFrame(X_pd, schema=schema)
del X_pd
data_model.printSchema()

root
|-- n_pregnant: integer (nullable = true)
|-- plasma_glucose_concentration: float (nullable = true)
|-- blood_pressure: float (nullable = true)
|-- triceps_skinfold_thickness: float (nullable = true)
|-- 2hours_serum_insulin: float (nullable = true)
|-- body_mass_index: float (nullable = true)
|-- diabetes_pedigree_function: float (nullable = true)
|-- age: integer (nullable = true)
|-- class: integer (nullable = true)
```

**Colunas presentes no Dataframe**

```
In [8]: for column in data_.limit(5).toPandas().columns:  
        print(column)
```

```
n_pregnant  
plasma_glucose_concentration  
blood_pressure  
triceps_skinfold_thickness  
2hours_serum_insulin  
body_mass_index  
diabetes_pedigree_function  
age  
class
```

*n\_pregnant - Número de vezes que engravidou*

*plasma\_glucose\_concentration - Concentração de glicose plasmática a 2 horas em um teste oral de tolerância à glicose*

*blood\_pressure - Pressão arterial diastólica (mm Hg)*

*triceps\_skinfold\_thickness - Espessura da dobra da pele do tríceps (mm)*

*2hours\_serum\_insulin - Insulina sérica de 2 horas (mu U / ml)*

*body\_mass\_index - Índice de massa corporal (peso em kg / (altura em m) ^ 2)*

*diabetes\_pedigree\_function - valores que medem a tendência ao desenvolvimento de diabetes com base nas relações genéticas do indivíduo*

*age - Idade*

*class - Presença ou não de diabetes*

**Vamos ver as estatísticas das variáveis**

```
In [9]: data_.summary().toPandas().T
```

```
Out[9]:
```

		0	1	2	3	4	
	summary	count	mean	stddev	min	25%	50%
	n_pregnant	768	3.8450520833333335	3.36957806269887	0	1	
plasma_glucose_concentration		768	120.89453125	31.97261819513622	0.0	99.0	117
blood_pressure		768	69.10546875	19.355807170644777	0.0	62.0	72
triceps_skinfold_thickness		768	20.536458333333332	15.952217567727642	0.0	0.0	23
2hours_serum_insulin		768	79.79947916666667	115.24400235133803	0.0	0.0	29
body_mass_index		768	31.99257813890775	7.884160293010772	0.0	27.3	32
diabetes_pedigree_function		768	0.4718763029280429	0.3313285967924436	0.078	0.243	0.3
age		768	33.240885416666664	11.760231540678689	21	24	
class		768	0.3489583333333333	0.476951377242799	0	0	


**Vamos plotar Boxplots das variáveis para ter uma visão geral de dimensão e distribuição dos dados**

```
In [58]: import matplotlib.pyplot as plt
import seaborn as sns

plt.rcParams['figure.figsize']=(20,10)
tips = sns.load_dataset("tips")

sns.boxplot(x="variable", y="value", data=pd.melt(data_.toPandas().drop(['class'], axis=1)))

sns.despine(offset=10, trim=True)
```



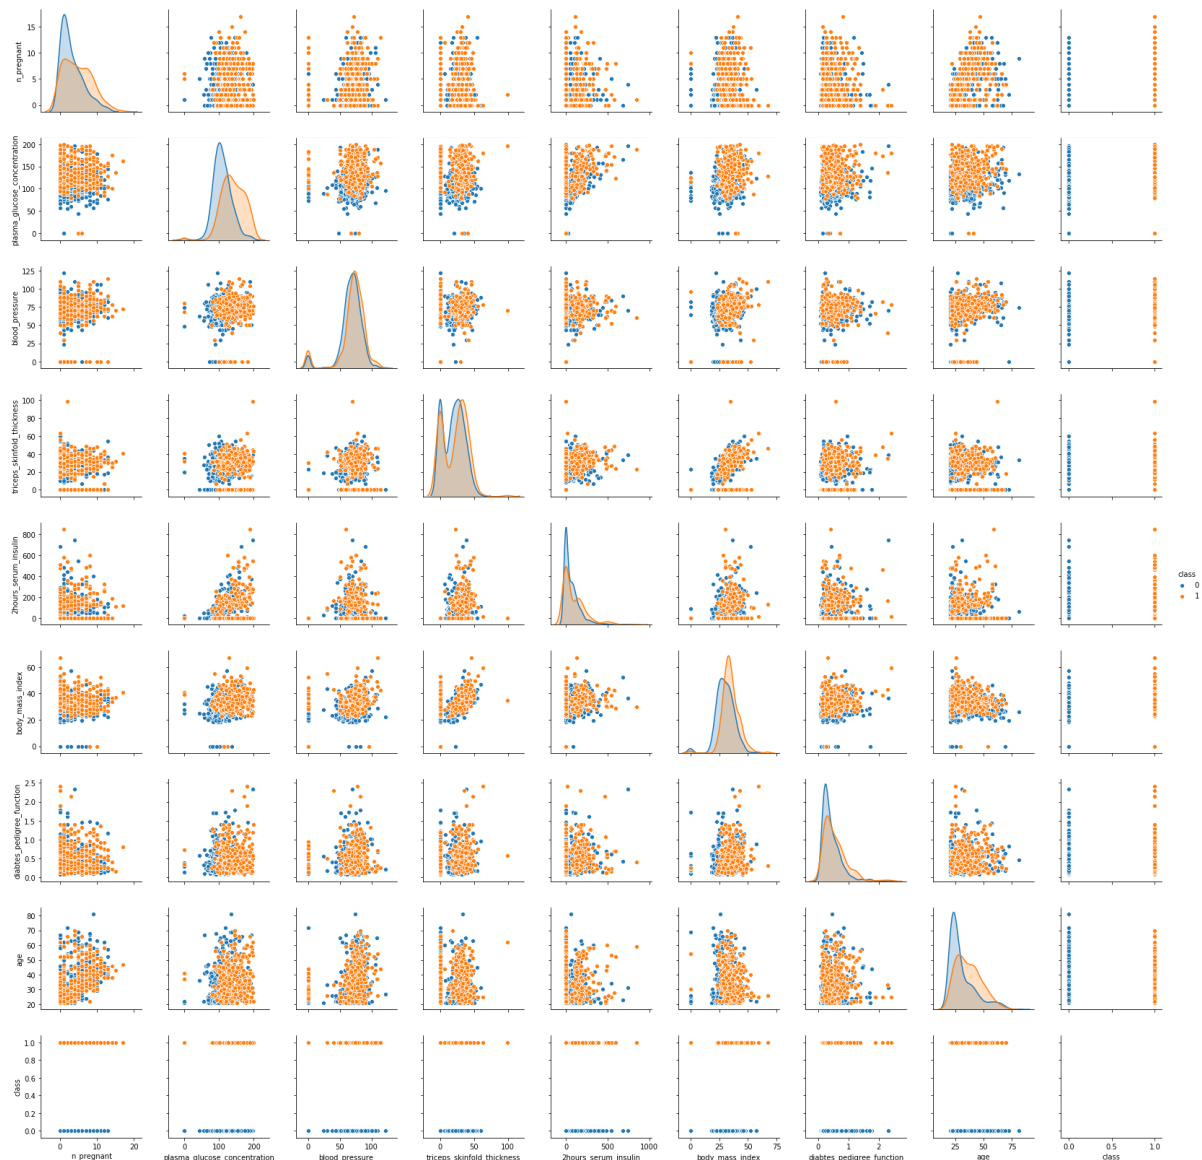
*No gráfico acima fica evidente que algumas variáveis apresentam bastantes outliers, o que pode impactar na performance dos modelos de Machine learning*

*Futuramente iremos tratá-los a fim de verificar o quanto esses outliers interferem no modelo*

***Vamos plotar uma Scatterplot Matrix para ver se conseguimos enxergar relações entre as variáveis, assim como se existe alguma Decision Boundery aparente aos nossos olhos***

```
In [11]: sns.pairplot(data_.toPandas(), hue="class");
```

```
/home/matheus/anaconda3/lib/python3.7/site-packages/statsmodels/nonparametric/kde.py:487: RuntimeWarning: invalid value encountered in true_divide
  binned = fast_linbin(X, a, b, gridsize) / (delta * nobs)
/home/matheus/anaconda3/lib/python3.7/site-packages/statsmodels/nonparametric/kdetools.py:34: RuntimeWarning: invalid value encountered in double_scalars
  FAC1 = 2*(np.pi*bw/RANGE)**2
```




Notas: > > >

**Agora que já tivemos um insight visual, vamos mensurar numericamente o qual forte são as relações entre essas variáveis usando a matrix de correlação**

```
In [59]: import numpy as np

data_corr = data_.toPandas().drop(['class'], axis=1).corr()

mask = np.triu(data_corr)
sns.heatmap(data_corr, annot=True, mask=mask, annot_kws={'rotation': 90});
```



In [ ]:

**Dentre as variáveis independentes do modelo "n\_pregnant" é a única que aparenta ter um comportamento discreto. Abaixo iremos explorar ela um pouco**


**Valores únicos e sua ocorrência:**

```
In [13]: data_.groupby('n_pregnant').count().toPandas().sort_values(by='count', ascending=False)
```

Out[13]:

	n_pregnant	count
1	1	135
16	0	111
15	2	103
4	3	75
9	4	68
5	5	57
3	6	50
11	7	45
10	8	38
7	9	28
12	10	24
13	11	11
2	13	10
0	12	9
14	14	2
6	15	1
8	17	1

```
In [60]: sns.distplot(data_.toPandas()["n_pregnant"]);
```



In [ ]:



Vamos fazer um teste de normalidade

```
In [15]: from scipy import stats

x = data_.select('n_pregnant').toPandas()['n_pregnant']

k2, p = stats.normaltest(x)
alpha = 1e-3
print("p = {:g}".format(p))

if p < alpha:
    print("The null hypothesis can be rejected")
else:
    print("The null hypothesis cannot be rejected")
```

```
p = 3.91429e-18
The null hypothesis can be rejected
```

O teste rejeitou a hipótese nula, consequentemente a coluna "n\_pregnant" não apresenta uma curva normal

## Criando os modelos ML

**Antes de começar o processamento do Dataframe vamos treinar um modelo sem e verificar as métricas com os dados crus a fim de criar um benchmark**

```
In [16]: stages = []

label_stringIdx = StringIndexer(inputCol = 'class', outputCol = 'label')
stages += [label_stringIdx]
assemblerInputs = ['n_pregnant', 'age', 'plasma_glucose_concentration', 'blood_
pressure', 'triceps_skinfold_thickness',
                    '2hours_serum_insulin', 'body_mass_index', 'diabetes_pedigree_functi
on']

assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
stages += [assembler]
```

```
In [17]: cols = data_.columns

pipeline = Pipeline(stages = stages)
pipelineModel = pipeline.fit(data_)
data_ = pipelineModel.transform(data_)
selectedCols = ['Label', 'features'] + cols
data_ = data_.select(selectedCols)
data_.printSchema()

root
|-- label: double (nullable = false)
|-- features: vector (nullable = true)
|-- n_pregnant: integer (nullable = true)
|-- plasma_glucose_concentration: float (nullable = true)
|-- blood_pressure: float (nullable = true)
|-- triceps_skinfold_thickness: float (nullable = true)
|-- 2hours_serum_insulin: float (nullable = true)
|-- body_mass_index: float (nullable = true)
|-- diabetes_pedigree_function: float (nullable = true)
|-- age: integer (nullable = true)
|-- class: integer (nullable = true)
```

```
In [18]: # Preparando pros modelos
train, test = data_.randomSplit([0.7, 0.3], seed = 2019)
print("Training Dataset Count: " + str(train.count()))
print("Test Dataset Count: " + str(test.count()))
```


Training Dataset Count: 544

Test Dataset Count: 224

## Logistic Regression Model


```
In [62]: from pyspark.ml.classification import LogisticRegression
Lr = LogisticRegression(featuresCol = 'features', labelCol = 'label', maxIter=
10)
LrModel = Lr.fit(train)

trainingSummary = LrModel.summary
roc = trainingSummary.roc.toPandas()
plt.plot(roc['FPR'],roc['TPR'])
plt.ylabel('False Positive Rate')
plt.xlabel('True Positive Rate')
plt.title('ROC Curve')
plt.show()
print('Training set areaUnderROC: ' + str(trainingSummary.areaUnderROC))
```



Training set areaUnderROC: 0.8412677324594429

```
In [63]: pr = trainingSummary.pr.toPandas()
plt.plot(pr['recall'],pr['precision'])
plt.ylabel('Precision')
plt.xlabel('Recall')
plt.show()
```



```
In [21]: from pyspark.ml.evaluation import BinaryClassificationEvaluator

predictions = LrModel.transform(test)
evaluator = BinaryClassificationEvaluator()
print('Test Area Under ROC', evaluator.evaluate(predictions))

Test Area Under ROC 0.8171812080536911
```

## Decision Tree Classifier

```
In [22]: from pyspark.ml.classification import DecisionTreeClassifier

dt = DecisionTreeClassifier(featuresCol = 'features', labelCol = 'label', maxDepth = 3)
dtModel = dt.fit(train)
predictions = dtModel.transform(test)
evaluator = BinaryClassificationEvaluator()
print("Test Area Under ROC: " + str(evaluator.evaluate(predictions, {evaluator.metricName: "areaUnderROC"})))

Test Area Under ROC: 0.720268456375839
```

## Random Forest Classifier

```
In [23]: from pyspark.ml.classification import RandomForestClassifier

rf = RandomForestClassifier(featuresCol = 'features', labelCol = 'label')
rfModel = rf.fit(train)
predictions = rfModel.transform(test)
evaluator = BinaryClassificationEvaluator()
print("Test Area Under ROC: " + str(evaluator.evaluate(predictions, {evaluator
.metricName: "areaUnderROC"})))
```

Test Area Under ROC: 0.8136912751677854

## Gradient-Boosted Tree Classifier

```
In [24]: from pyspark.ml.classification import GBTClassifier

gbt = GBTClassifier(maxIter=10)
gbtModel = gbt.fit(train)
predictions = gbtModel.transform(test)
evaluator = BinaryClassificationEvaluator()
print("Test Area Under ROC: " + str(evaluator.evaluate(predictions, {evaluator
.metricName: "areaUnderROC"})))
```

Test Area Under ROC: 0.7434004474272928

## Gradient-Boosted Tree Classifier + Cross Validator + Grid Search

```
In [25]: from pyspark.ml.tuning import ParamGridBuilder, CrossValidator

paramGrid = (ParamGridBuilder()
              .addGrid(gbt.maxDepth, [2, 4, 6])
              .addGrid(gbt.maxBins, [20, 60])
              .addGrid(gbt.maxIter, [10, 20])
              .build())

cv = CrossValidator(estimator=gbt, estimatorParamMaps=paramGrid, evaluator=eva
luator, numFolds=5)
# Run cross validations. This can take about 6 minutes since it is training o
ver 20 trees!
cvModel = cv.fit(train)
predictions = cvModel.transform(test)
evaluator.evaluate(predictions)
```

Out[25]: 0.8207606263982102

## Nosso Benchmark

*Logistic Regression Model - Area Under ROC 0.8171812080536911*

*Decision Tree Classifier - Area Under ROC: 0.720268456375839*

*Random Forest Classifier - Area Under ROC: 0.8136912751677854*

*Gradient-Boosted Tree Classifier - Area Under ROC: 0.7434004474272928*

*Gradient-Boosted Tree Classifier + Cross Validator + Grid Search - Area Under ROC:0.8207606263982102*

## Criando nosso modelo

**Vamos desenvolver algumas funções de processamento, treinamento e teste, para que em seguida possamos rodar todas possibilidades de modelos, considerando varios tipos de processamento diferentes**

Funções de processamento:

```
In [26]: from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.classification import GBTClassifier
from pyspark.ml.classification import DecisionTreeClassifier
from pyspark.ml.classification import MultilayerPerceptronClassifier
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
from pyspark.ml.evaluation import BinaryClassificationEvaluator
from pyspark.mllib.evaluation import BinaryClassificationMetrics
from pyspark.ml.classification import LogisticRegression
from pyspark.ml.feature import MinMaxScaler, StandardScaler
from pyspark.ml import Pipeline
from pyspark.sql import functions as f
```

Remover outliers

```
In [44]: # outliers_parameters = [False, 'StdDev']

def remove_outliers(df, outliers_type):
    if outliers_type == 'StdDev':
        data = df.where(df.body_mass_index > 10.0)
        data = data.where(data.triceps_skinfold_thickness > 0)
        data = data.where(data['2hours_serum_insulin'] > 0)
        data = data.where(data['2hours_serum_insulin'] < 350)
        return data

    else:
        return df
```

### Fazer amostragem nas classes - balanceamento

```
In [28]: # sample_parameters = [False, True]

def balance_dataframe(df, sample_type):
    if sample_type == True:
        lenght = []
        sample_list = []
        lenght.append(df.filter(df['class'] == 1).count())
        lenght.append(df.filter(df['class'] == 0).count())
        sample_size = min(lenght)
        sample_list.extend(df.filter(df['class'] == 1).rdd.takeSample(False, sample_size, seed=0))
        sample_list.extend(df.filter(df['class'] == 0).rdd.takeSample(False, sample_size, seed=0))

        data = sqlContext.createDataFrame(sample_list)
        return data
    else:
        return df
```

### Remover agrupamento de duplicados remanecentes

```
In [29]: # remove_distinct_parameters = [False, ['n_pregnant', '2hours_serum_insulin', 'age', 'class'], ['age', 'class']]

def remove_distinct(df, remove_distinct_type):
    if remove_distinct_type == False:
        return df
    else:
        data = df.drop_duplicates(subset=remove_distinct_type)
        return data
```

*Normalizar valores - MinMax / Standalizer*



```

In [30]: # normalization_parameters = [False, 'MinMax', 'Standalizer']

def normalize_dataframe(df, normalization_type):

    columns_to_scale = []
    stages = []

    for column in df.columns:
        if column == "class": continue
        columns_to_scale.append(column)

    if normalization_type == 'MinMax':

        label_stringIdx = StringIndexer(inputCol = 'class', outputCol = 'label')
        stages += [label_stringIdx]
        pipeline = Pipeline(stages = stages)
        pipelineModel = pipeline.fit(df)
        df_result = pipelineModel.transform(df)

        assemblers = VectorAssembler(inputCols=columns_to_scale, outputCol="features_preprocessing")
        vector = assemblers.transform(df_result)
        standard_scaler = MinMaxScaler(inputCol="features_preprocessing", outputCol="features")
        train = standard_scaler.fit(vector).transform(vector)
        return train

    elif normalization_type == 'Standalizer':

        label_stringIdx = StringIndexer(inputCol = 'class', outputCol = 'label')
        stages += [label_stringIdx]
        pipeline = Pipeline(stages = stages)
        pipelineModel = pipeline.fit(df)
        df_result = pipelineModel.transform(df)

        assemblers = VectorAssembler(inputCols=columns_to_scale, outputCol="features_preprocessing")
        vector = assemblers.transform(df_result)
        standard_scaler = StandardScaler(inputCol="features_preprocessing", outputCol="features")
        train = standard_scaler.fit(vector).transform(vector)
        return train

    else:

        label_stringIdx = StringIndexer(inputCol = 'class', outputCol = 'label')
        stages += [label_stringIdx]

        assembler = VectorAssembler(inputCols=columns_to_scale, outputCol="features")

```

```

stages += [assembler]

cols = df.columns

pipeline = Pipeline(stages = stages)
pipelineModel = pipeline.fit(df)
df_result = pipelineModel.transform(df)
selectedCols = ['Label', 'features'] + cols
df_result = df_result.select(selectedCols)
return df_result

```

Modelos:

```

In [46]: def apply_Logistic_Regression_Model(df):

    train, test = df.randomSplit([0.8, 0.2], seed = 2019)

    #tocar essa linha
    model = LogisticRegression(featuresCol = 'features', labelCol = 'Label', maxIter=10)

    evaluator = BinaryClassificationEvaluator()

    # Create ParamGrid for Cross Validation
    paramGrid = (ParamGridBuilder()
                 .addGrid(model.regParam,[0.02])
                 .addGrid(model.elasticNetParam,[0.2])
                 .build())

    # Create 5-fold CrossValidator
    cv = CrossValidator(estimator = model,
                        estimatorParamMaps = paramGrid,
                        evaluator = evaluator,
                        numFolds = 5)

    # Run cross validations.
    cvModel = cv.fit(train)

    # Use test set here so we can measure the accuracy of our model on new data
    predictions = cvModel.transform(test)

    # cvModel uses the best model found from the Cross Validation
    # Evaluate best model
    acc = evaluator.evaluate(predictions)
    auc = BinaryClassificationMetrics(predictions['Label', 'prediction'].rdd).areaUnderROC
    print('Accuracy:', acc)
    print('AUC:', auc)
    return acc, auc

```

```

In [47]: def apply_Decision_Tree_Model(df):
    train, test = df.randomSplit([0.8, 0.2], seed = 2019)

    #tocar essa linha
    model = DecisionTreeClassifier(featuresCol = 'features', labelCol = 'Label')

    evaluator = BinaryClassificationEvaluator()

    # Create ParamGrid for Cross Validation
    paramGrid = (ParamGridBuilder()
        .addGrid(model.maxDepth, [2, 20])
        .addGrid(model.maxBins, [10, 40])
        .build())

    # Create 5-fold CrossValidator
    cv = CrossValidator(estimator = model,
        estimatorParamMaps = paramGrid,
        evaluator = evaluator,
        numFolds = 5)

    # Run cross validations.
    cvModel = cv.fit(train)

    # Use test set here so we can measure the accuracy of our model on new data
    predictions = cvModel.transform(test)

    # cvModel uses the best model found from the Cross Validation
    # Evaluate best model
    acc = evaluator.evaluate(predictions)
    auc = BinaryClassificationMetrics(predictions['Label', 'prediction'].rdd).areaUnderROC
    print('Accuracy:', acc)
    print('AUC:', auc)
    return acc, auc

```

```
In [48]: def apply_Random_Forest_Model(df):

    train, test = df.randomSplit([0.8, 0.2], seed = 2019)
    # Create an initial RandomForest model.
    model = RandomForestClassifier(labelCol="Label", featuresCol="features")
    # Evaluate model
    evaluator = BinaryClassificationEvaluator()
    # Create ParamGrid for Cross Validation
    paramGrid = (ParamGridBuilder()
                 .addGrid(model.maxDepth, [2 ])
                 .addGrid(model.maxBins, [10,])
                 .addGrid(model.numTrees, [5])
                 .build())

    # Create 5-fold CrossValidator
    cv = CrossValidator(estimator = model,
                       estimatorParamMaps = paramGrid,
                       evaluator = evaluator,
                       numFolds = 5)

    # Run cross validations.
    cvModel = cv.fit(train)

    # Use test set here so we can measure the accuracy of our model on new data
    predictions = cvModel.transform(test)

    # cvModel uses the best model found from the Cross Validation
    # Evaluate best model
    acc = evaluator.evaluate(predictions)
    auc = BinaryClassificationMetrics(predictions['Label', 'prediction'].rdd).areaUnderROC

    print('Accuracy:', acc)
    print('AUC:', auc)
    return acc, auc
```

```

In [49]: def apply_Gradient_BoostedT_Model(df):

    train, test = df.randomSplit([0.8, 0.2], seed = 2019)

    #tocar essa linha
    model = GBTClassifier(featuresCol = 'features', labelCol = 'label', maxIter=10 )

    evaluator = BinaryClassificationEvaluator()

    # Create ParamGrid for Cross Validation
    paramGrid = (ParamGridBuilder()
        .addGrid(model.maxDepth, [2, ])
        .addGrid(model.maxBins, [20])
        .addGrid(model.maxIter, [10])
        .build())

    # Create 5-fold CrossValidator
    cv = CrossValidator(estimator = model,
        estimatorParamMaps = paramGrid,
        evaluator = evaluator,
        numFolds = 5)

    # Run cross validations.
    cvModel = cv.fit(train)

    # Use test set here so we can measure the accuracy of our model on new data
    predictions = cvModel.transform(test)

    # cvModel uses the best model found from the Cross Validation
    # Evaluate best model
    acc = evaluator.evaluate(predictions)
    auc = BinaryClassificationMetrics(predictions['label', 'prediction'].rdd).areaUnderROC
    print('Accuracy:', acc)
    print('AUC:', auc)
    return acc, auc

```

```
In [50]: def apply_MultiLayer_Perceptron(df):
    train, test = df.randomSplit([0.8, 0.2], seed = 2019)

    model = MultilayerPerceptronClassifier(featuresCol = 'features', labelCol
    = 'Label')
    evaluator = BinaryClassificationEvaluator()
    paramGrid = (ParamGridBuilder()
        .addGrid(model.maxIter, [10, 20])
        .addGrid(model.stepSize, [0.03, 0.01])
        .addGrid(model.solver, ['L-bfgs', 'gd'])
        .addGrid(model.layers, [[8, 12, 2], [8, 5, 2], [8, 5, 4, 2]])
        .build())

    # Create 5-fold CrossValidator
    cv = CrossValidator(estimator = model, estimatorParamMaps = paramGrid, evaluator = evaluator, numFolds = 5)
    # Run cross validations.
    cvModel = cv.fit(train)

    # Use test set here so we can measure the accuracy of our model on new data
    predictions = cvModel.transform(test)

    # cvModel uses the best model found from the Cross Validation
    # Evaluate best model
    acc = f
    evaluator.evaluate(predictions)
    auc = BinaryClassificationMetrics(predictions['Label', 'prediction'].rdd).areaUnderROC
    print('Accuracy:', acc)
    print('AUC:', auc)
    return acc, auc
```

### Ordem que as funções devem rodar

```
data = removeoutliers(data, outliers_tag)

data = balance_dataframe(data, sample_tag)

data = remove_distinct(data, remove_distinct_tag)

data = normalize_dataframe(data, normalization_tag)

apply_Logistic_Regression_Model(data)
```

**No código a baixo iremos fazer um loop entre todos métodos de processamento de modelos de machine learnin para encontrar qual o melhor pipeline de processamento/modelo para nosso conjunto de dados**

```

In [51]: from datetime import datetime
start = datetime.now()

# models = ["MLP", "GBT", "RDF", "DCT", "LGR"]

outliers_parameters = [False, 'StdDev', 'DropCol']
sample_parameters = [False, True]
remove_distinct_parameters = [False, ['n_pregnant', '2hours_serum_insulin', 'age', 'class'], ['age', 'class']]
normalization_parameters = [False, 'MinMax', 'Standalizer']
models = ["RDF", "LGR", "GBT"]

result = []

for model in models:
    for outliers_tag in outliers_parameters:
        for sample_tag in sample_parameters:
            for remove_distinct_tag in remove_distinct_parameters:
                for normalization_tag in normalization_parameters:

                    dic = {}

                    data = remove_outliers(data_model, outliers_tag)
                    data = balance_dataframe(data, sample_tag)
                    data = remove_dinstinct(data, remove_distinct_tag)
                    data = normalize_dataframe(data, normalization_tag)

                    print(model)
                    if model == "MLP":
                        acc, auc = apply_MultiLayer_Perceptron(data)

                    elif model == "GBT":
                        acc, auc = apply_Gradient_BoostedT_Model(data)

                    elif model == "RDF":
                        acc, auc = apply_Random_Forest_Model(data)

                    elif model == "DCT":
                        acc, auc = apply_Decision_Tree_Model(data)

                    elif model == "LGR":
                        acc, auc = apply_Logistic_Regression_Model(data)

                    print(" ")
                    dic["Model Name"] = model
                    dic["outliers_tag"] = outliers_tag
                    dic["sample_tag"] = sample_tag
                    dic["remove_distinct_tag"] = remove_distinct_tag
                    dic["normalization_tag"] = normalization_tag
                    dic["Acuracia"] = acc
                    dic["AUC"] = auc
                    result.append(dic)

df = pd.DataFrame(result)

```

```
df = df.sort_values(by=['AUC'], ascending=False)

stop = datetime.now()
print('Time: ', stop - start)
```



RDF

Accuracy: 0.8453608247422681

AUC: 0.7697368421052632

RDF

Accuracy: 0.7912234042553192

AUC: 0.6758209822725951

RDF

Accuracy: 0.7912234042553192

AUC: 0.6758209822725951

RDF

Accuracy: 0.7703488372093023

AUC: 0.7881578947368422

RDF

Accuracy: 0.8031666666666665

AUC: 0.8331043956043955

RDF

Accuracy: 0.8031666666666665

AUC: 0.8331043956043955

RDF

Accuracy: 0.6212121212121212

AUC: 0.5833333333333334

RDF

Accuracy: 0.6262626262626262

AUC: 0.5625

RDF

Accuracy: 0.6262626262626262

AUC: 0.5625

RDF

Accuracy: 0.8359113712374582

AUC: 0.7317124735729387

RDF

Accuracy: 0.8292224080267557

AUC: 0.7609427609427609

RDF

Accuracy: 0.8292224080267557

AUC: 0.7609427609427609

RDF

Accuracy: 0.7361793611793612

AUC: 0.6512422360248447

RDF

Accuracy: 0.7751842751842751

AUC: 0.7111111111111111

RDF

Accuracy: 0.7751842751842751  
AUC: 0.7111111111111111

RDF  
Accuracy: 0.8  
AUC: 0.775

RDF  
Accuracy: 0.9861111111111112  
AUC: 0.875

RDF  
Accuracy: 0.9861111111111112  
AUC: 0.875

RDF  
Accuracy: 0.8771991555242786  
AUC: 0.7384615384615384

RDF  
Accuracy: 0.8496428571428571  
AUC: 0.8082706766917294

RDF  
Accuracy: 0.8496428571428571  
AUC: 0.8082706766917294

RDF  
Accuracy: 0.8968750000000001  
AUC: 0.8875

RDF  
Accuracy: 0.950625  
AUC: 0.962962962962963

RDF  
Accuracy: 0.950625  
AUC: 0.962962962962963

RDF  
Accuracy: 0.44696969696969696  
AUC: 0.38636363636363635

RDF  
Accuracy: 0.5857142857142856  
AUC: 0.5347222222222222

RDF  
Accuracy: 0.5857142857142856  
AUC: 0.5347222222222222

RDF  
Accuracy: 0.7391304347826086  
AUC: 0.735

RDF  
Accuracy: 0.8063241106719368

AUC: 0.7685185185185186

RDF

Accuracy: 0.8063241106719368

AUC: 0.7685185185185186

RDF

Accuracy: 0.8527667984189723

AUC: 0.7559288537549407

RDF

Accuracy: 0.8241106719367588

AUC: 0.8003952569169962

RDF

Accuracy: 0.8241106719367588

AUC: 0.8003952569169962

RDF

Accuracy: 0.8

AUC: 0.8125

RDF

Accuracy: 0.6666666666666666

AUC: 0.7857142857142857

RDF

Accuracy: 0.6666666666666666

AUC: 0.7857142857142857

RDF

Accuracy: 0.8453608247422681

AUC: 0.7697368421052632

RDF

Accuracy: 0.7912234042553192

AUC: 0.6758209822725951

RDF

Accuracy: 0.7912234042553192

AUC: 0.6758209822725951

RDF

Accuracy: 0.7703488372093023

AUC: 0.7881578947368422

RDF

Accuracy: 0.8031666666666665

AUC: 0.8331043956043955

RDF

Accuracy: 0.8031666666666665

AUC: 0.8331043956043955

RDF

Accuracy: 0.6212121212121212

AUC: 0.5833333333333334

RDF

Accuracy: 0.6262626262626262

AUC: 0.5625

RDF

Accuracy: 0.6262626262626262

AUC: 0.5625

RDF

Accuracy: 0.8359113712374582

AUC: 0.7317124735729387

RDF

Accuracy: 0.8292224080267557

AUC: 0.7609427609427609

RDF

Accuracy: 0.8292224080267557

AUC: 0.7609427609427609

RDF

Accuracy: 0.7361793611793612

AUC: 0.6512422360248447

RDF

Accuracy: 0.7751842751842751

AUC: 0.7111111111111111

RDF

Accuracy: 0.7751842751842751

AUC: 0.7111111111111111

RDF

Accuracy: 0.8

AUC: 0.775

RDF

Accuracy: 0.9861111111111112

AUC: 0.875

RDF

Accuracy: 0.9861111111111112

AUC: 0.875

LGR

Accuracy: 0.7993127147766323

AUC: 0.7573049987925623

LGR

Accuracy: 0.7429078014184399

AUC: 0.6794871794871794

LGR

Accuracy: 0.6914893617021277

AUC: 0.6462201591511936

LGR

Accuracy: 0.721576227390181

AUC: 0.8214285714285714

LGR

Accuracy: 0.7693333333333339

AUC: 0.717391304347826

LGR

Accuracy: 0.7403333333333332

AUC: 0.7083333333333334

LGR

Accuracy: 0.5858585858585857

AUC: 0.6499999999999999

LGR

Accuracy: 0.6565656565656565

AUC: 0.5625

LGR

Accuracy: 0.6363636363636364

AUC: 0.6499999999999999

LGR

Accuracy: 0.7391304347826085

AUC: 0.6599468320779797

LGR

Accuracy: 0.8189799331103679

AUC: 0.7541806020066888

LGR

Accuracy: 0.774247491638796

AUC: 0.7234800838574424

LGR

Accuracy: 0.6480343980343977

AUC: 0.5785984848484848

LGR

Accuracy: 0.6689189189189189

AUC: 0.6328029375764993

LGR

Accuracy: 0.6744471744471747

AUC: 0.5872015915119363

LGR

Accuracy: 0.5714285714285714

AUC: 0.4318181818181819

LGR

Accuracy: 0.6805555555555556

AUC: 0.5357142857142857

LGR

Accuracy: 0.7361111111111113  
AUC: 0.5833333333333334

LGR  
Accuracy: 0.7424349049964815  
AUC: 0.6825396825396826

LGR  
Accuracy: 0.7864285714285715  
AUC: 0.7564935064935066

LGR  
Accuracy: 0.7228571428571428  
AUC: 0.7000000000000001

LGR  
Accuracy: 0.7962500000000005  
AUC: 0.7639225181598064

LGR  
Accuracy: 0.89  
AUC: 0.9166666666666667

LGR  
Accuracy: 0.8700000000000003  
AUC: 0.825

LGR  
Accuracy: 0.7575757575757576  
AUC: 0.6083333333333334

LGR  
Accuracy: 0.8999999999999999  
AUC: 0.7285714285714286

LGR  
Accuracy: 0.8428571428571427  
AUC: 0.7285714285714286

LGR  
Accuracy: 0.6996047430830039  
AUC: 0.7150000000000001

LGR  
Accuracy: 0.7312252964426879  
AUC: 0.7601214574898786

LGR  
Accuracy: 0.7154150197628459  
AUC: 0.6887351778656127

LGR  
Accuracy: 0.7272727272727273  
AUC: 0.66699604743083

LGR  
Accuracy: 0.6818181818181819

AUC: 0.67

LGR

Accuracy: 0.6897233201581028

AUC: 0.67

LGR

Accuracy: 0.48

AUC: 0.2222222222222222

LGR

Accuracy: 0.625

AUC: 0.75

LGR

Accuracy: 0.45833333333333326

AUC: 0.0

LGR

Accuracy: 0.7993127147766323

AUC: 0.7573049987925623

LGR

Accuracy: 0.7429078014184399

AUC: 0.6794871794871794

LGR

Accuracy: 0.6914893617021277

AUC: 0.6462201591511936

LGR

Accuracy: 0.721576227390181

AUC: 0.8214285714285714

LGR

Accuracy: 0.7693333333333339

AUC: 0.717391304347826

LGR

Accuracy: 0.7403333333333332

AUC: 0.7083333333333334

LGR

Accuracy: 0.5858585858585857

AUC: 0.6499999999999999

LGR

Accuracy: 0.6565656565656565

AUC: 0.5625

LGR

Accuracy: 0.6363636363636364

AUC: 0.6499999999999999

LGR

Accuracy: 0.7391304347826085

AUC: 0.6599468320779797

LGR

Accuracy: 0.8189799331103679

AUC: 0.7541806020066888

LGR

Accuracy: 0.774247491638796

AUC: 0.7234800838574424

LGR

Accuracy: 0.6480343980343977

AUC: 0.5785984848484848

LGR

Accuracy: 0.6689189189189189

AUC: 0.6328029375764993

LGR

Accuracy: 0.6744471744471747

AUC: 0.5872015915119363

LGR

Accuracy: 0.5714285714285714

AUC: 0.4318181818181819

LGR

Accuracy: 0.6805555555555556

AUC: 0.5357142857142857

LGR

Accuracy: 0.7361111111111113

AUC: 0.5833333333333334

GBT

Accuracy: 0.8324169530355098

AUC: 0.7327127659574468

GBT

Accuracy: 0.8262411347517731

AUC: 0.6997549019607843

GBT

Accuracy: 0.8257978723404256

AUC: 0.6997549019607843

GBT

Accuracy: 0.8338178294573643

AUC: 0.7282608695652174

GBT

Accuracy: 0.8443333333333333

AUC: 0.8219533275713051

GBT

Accuracy: 0.8443333333333333

AUC: 0.8219533275713051



GBT  
Accuracy: 0.6565656565656566  
AUC: 0.696969696969697

GBT  
Accuracy: 0.7828282828282829  
AUC: 0.75

GBT  
Accuracy: 0.7828282828282829  
AUC: 0.75

GBT  
Accuracy: 0.8520066889632106  
AUC: 0.7670004171881519

GBT  
Accuracy: 0.8808528428093647  
AUC: 0.7874426366291197

GBT  
Accuracy: 0.8808528428093647  
AUC: 0.7874426366291197

GBT  
Accuracy: 0.7733415233415232  
AUC: 0.7361111111111112

GBT  
Accuracy: 0.7733415233415232  
AUC: 0.7361111111111112

GBT  
Accuracy: 0.7571428571428571  
AUC: 0.7571428571428572

GBT  
Accuracy: 0.9305555555555555  
AUC: 0.9090909090909091

GBT  
Accuracy: 0.9305555555555555  
AUC: 0.9090909090909091

GBT  
Accuracy: 0.8662913441238564  
AUC: 0.8712121212121212

GBT  
Accuracy: 0.8492857142857142  
AUC: 0.75

GBT  
Accuracy: 0.8485714285714286  
AUC: 0.75

GBT

Accuracy: 0.8706250000000001  
AUC: 0.8101903695408734

GBT  
Accuracy: 0.9012500000000004  
AUC: 0.8607843137254902

GBT  
Accuracy: 0.9012500000000004  
AUC: 0.8607843137254902

GBT  
Accuracy: 0.7272727272727273  
AUC: 0.6439393939393939

GBT  
Accuracy: 0.8785714285714286  
AUC: 0.7916666666666667

GBT  
Accuracy: 0.8785714285714286  
AUC: 0.7916666666666667

GBT  
Accuracy: 0.7885375494071146  
AUC: 0.7351190476190477

GBT  
Accuracy: 0.7658102766798419  
AUC: 0.6887351778656127

GBT  
Accuracy: 0.7658102766798419  
AUC: 0.6887351778656127

GBT  
Accuracy: 0.8448616600790514  
AUC: 0.7559523809523809

GBT  
Accuracy: 0.8043478260869567  
AUC: 0.7114624505928854

GBT  
Accuracy: 0.8043478260869567  
AUC: 0.7114624505928854

GBT  
Accuracy: 0.6799999999999999  
AUC: 0.8125

GBT  
Accuracy: 0.625  
AUC: 0.625

GBT  
Accuracy: 0.625

AUC: 0.625

GBT

Accuracy: 0.8324169530355098

AUC: 0.7327127659574468

GBT

Accuracy: 0.8262411347517731

AUC: 0.6997549019607843

GBT

Accuracy: 0.8257978723404256

AUC: 0.6997549019607843

GBT

Accuracy: 0.8338178294573643

AUC: 0.7282608695652174

GBT

Accuracy: 0.8443333333333333

AUC: 0.8219533275713051

GBT

Accuracy: 0.8443333333333333

AUC: 0.8219533275713051

GBT

Accuracy: 0.6565656565656566

AUC: 0.696969696969697

GBT

Accuracy: 0.7828282828282829

AUC: 0.75

GBT

Accuracy: 0.7828282828282829

AUC: 0.75

GBT

Accuracy: 0.8520066889632106

AUC: 0.7670004171881519

GBT

Accuracy: 0.8808528428093647

AUC: 0.7874426366291197

GBT

Accuracy: 0.8808528428093647

AUC: 0.7874426366291197

GBT

Accuracy: 0.7699631449631451

AUC: 0.7014742014742015

GBT

Accuracy: 0.7733415233415232

AUC: 0.7361111111111112

GBT

Accuracy: 0.7733415233415232

AUC: 0.7361111111111112

GBT

Accuracy: 0.7571428571428571

AUC: 0.7571428571428572

GBT

Accuracy: 0.9305555555555555

AUC: 0.9090909090909091

GBT

Accuracy: 0.9305555555555555

AUC: 0.9090909090909091

Time: 2:51:16.107462

In [54]: `df.head(15)`

Out[54]:

	Model Name	outliers_tag	sample_tag	remove_distinct_tag	normalization_tag	Acuracia	AUC
22	RDF	StdDev	False	[n_pregnant, 2hours_serum_insulin, age, class]	MinMax	0.950625	0.962961
23	RDF	StdDev	False	[n_pregnant, 2hours_serum_insulin, age, class]	Standalizer	0.950625	0.962961
76	LGR	StdDev	False	[n_pregnant, 2hours_serum_insulin, age, class]	MinMax	0.890000	0.916667
124	GBT	False	True	[age, class]	MinMax	0.930556	0.909091
125	GBT	False	True	[age, class]	Standalizer	0.930556	0.909091
161	GBT	DropCol	True	[age, class]	Standalizer	0.930556	0.909091
160	GBT	DropCol	True	[age, class]	MinMax	0.930556	0.909091
21	RDF	StdDev	False	[n_pregnant, 2hours_serum_insulin, age, class]	False	0.896875	0.887500
52	RDF	DropCol	True	[age, class]	MinMax	0.986111	0.875000
53	RDF	DropCol	True	[age, class]	Standalizer	0.986111	0.875000
16	RDF	False	True	[age, class]	MinMax	0.986111	0.875000
17	RDF	False	True	[age, class]	Standalizer	0.986111	0.875000
126	GBT	StdDev	False	False	False	0.866291	0.871212
131	GBT	StdDev	False	[n_pregnant, 2hours_serum_insulin, age, class]	Standalizer	0.901250	0.860784
130	GBT	StdDev	False	[n_pregnant, 2hours_serum_insulin, age, class]	MinMax	0.901250	0.860784